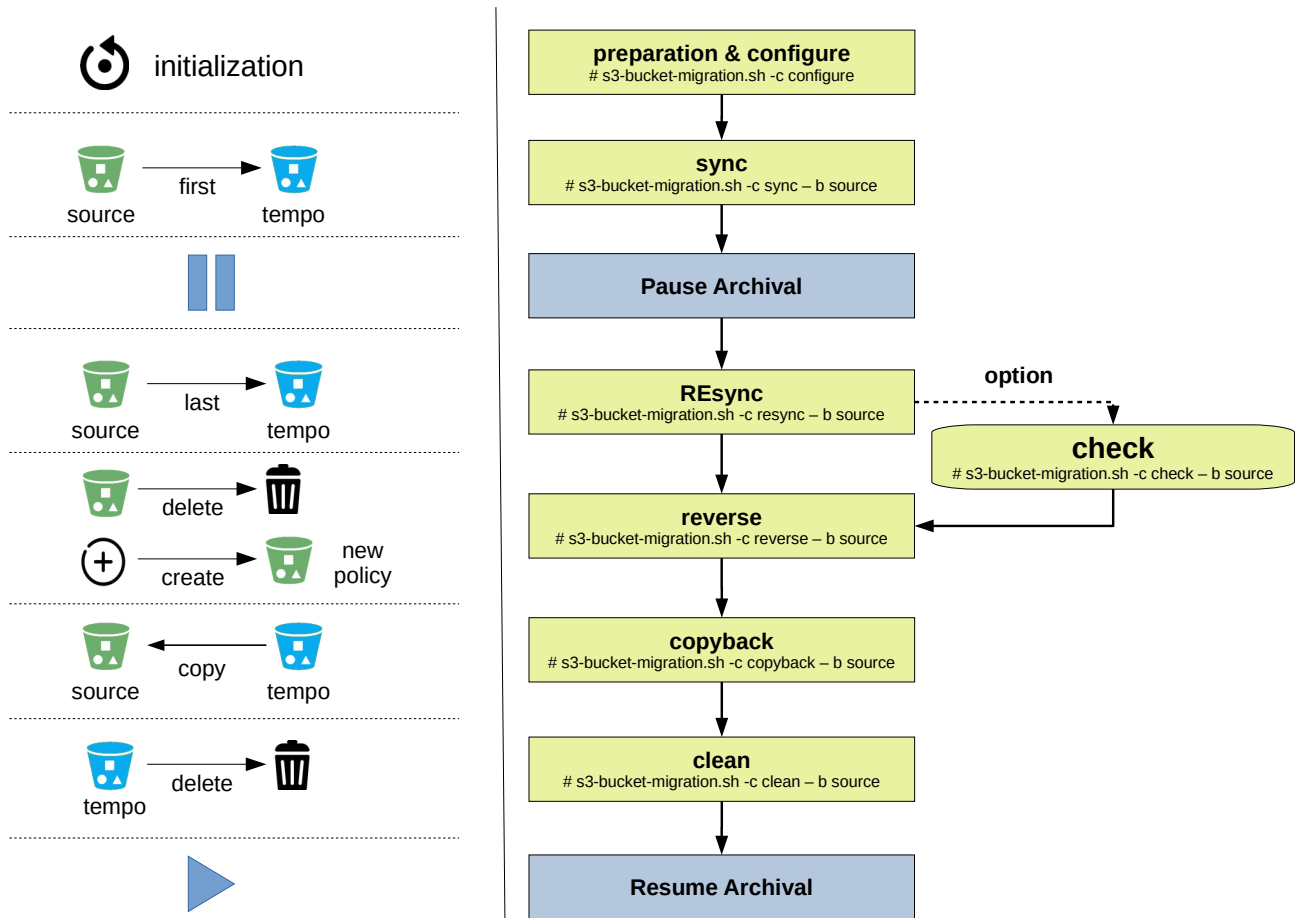


MIGRATION TOOL
for Rubrik + Cloudian buckets
v1.0b

1 Flowchart

The script “s3-bucket-migration.sh” follows the flowchart below :



2 Installation

You must download 2 files :

- the script : `s3-bucket-migration.sh`
- the package : `rclone-linux-amd64.rpm`

Then, proceed to the rpm installation :

```
[root@peter]# rpm -i ./rclone-linux-amd64.rpm
[root@peter]#
```

and run the script :

```
[root@peter]# ./s3-bucket-migration.sh -c configure
[root@peter]#
```

Notice : this script will use the Clouddian Admin API to purge the buckets, so you have to change the default password if needed.

3 Full example following the flowchart

This step is just to show you the syntax and all options available :

```
[root@peter]# ./s3-bucket-migration.sh
Error in the command line.
Use the format : ./s3-bucket-migration.sh -c <command> -b <bucketname>
<command> could be : configure , sync , resync , check , reverse , copyback , clean
```

Configuration of the s3 credentials and the s3 EndPoint :

```
[root@peter]# ./s3-bucket-migration.sh -c configure

Disclaimer / Warning
Use this tool with precautions (review the config file created manually for a double-check) for your environment :
it is NOT an official tool supported by Cloudian.
Cloudian can NOT be involved for any bugs or misconfiguration due to this tool. So you are using it at your own
risks and be aware of the restrictions.

--- Creating the configuration file with the AccessKey and SecretKey to manage the buckets ---
Please provide the AccessKey for the migration : 00f3bd5a0917d56db876
Please provide the SecretKey for the migration : TNFT6VEZHj6lB0Bfd8eOyufY3+t+4Ty1xC2ptv7d
Please provide the s3 EndPoint for the migration (example : http://s3-region.domainname) : http://s3-region.demo
Configuration done.
You can now proceed with the 'sync' command : ./s3-bucket-migration.sh -c sync -b <bucketname>
[root@peter]#
```

First sync in order to copy the maximum of objects before a full re-sync :

```
[root@peter]# ./s3-bucket-migration.sh -c sync -b plaineco-rubrik-0

Disclaimer / Warning
Use this tool with precautions (review the config file created manually for a double-check) for your environment :
it is NOT an official tool supported by Cloudian. Cloudian can NOT be involved for any bugs or misconfiguration
due to this tool. So you are using it at your own risks and be aware of the restrictions.

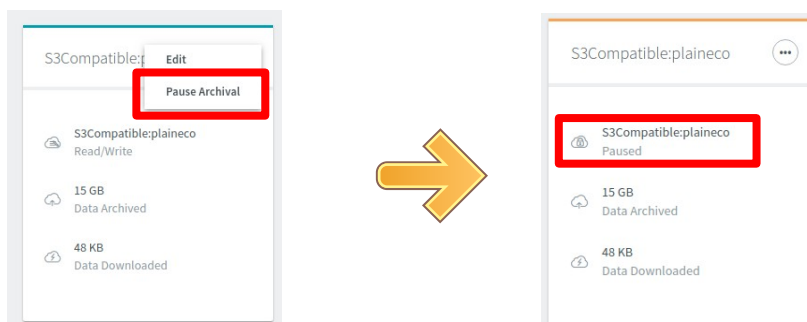
--- Preparing to migrate the bucket : plaineco-rubrik-0 ---
Do you want to continue (type : yes to continue, or anything else to abort) ? yes
OK, let's Go !
Checking the source bucket : plaineco-rubrik-0
    0 2020-01-29 18:47:11      -1 blobstore
    0 2020-01-29 18:47:11      -1 snappables

=> Bucket plaineco-rubrik-0 is available and readable. Continuing ...

Creating a new logfile : plaineco-rubrik-0.log in the current directory.
Creation of a temporary bucket to migrate firstly the data ... : plaineco-rubrik-0-tempo

Bucket plaineco-rubrik-0-tempo is now available. Continuing ...
Operation done.
--- Synchronization from the bucket plaineco-rubrik-0 to the bucket plaineco-rubrik-0-tempo : ---
Transferred:      13.841G / 13.841 GBytes, 100%, 47.289 MBytes/s, ETA 0s
Transferred:      92 / 92, 100%
Elapsed time:      4m59.7s
Please, change the ArchiveLocation parameter and force it in 'Pause mode'
Then, use the command re-sync : ./s3-bucket-migration.sh -c resync -b plaineco-rubrik-0
[root@peter]#
```

Then, you must go to the Rubrik GUI and force the “Pause Archival” mode on the ArchivalLocation before the full-resync (avoid upload of new objects during the full-resync) :



Full and last re-sync when no new uploads are performed by Rubrik (Pause mode) :

```
[root@peter]# ./s3-bucket-migration.sh -c resync -b plaineco-rubrik-0

Disclaimer / Warning
Use this tool with precautions (review the config file created manually for a double-check) for your environment :
it is NOT an official tool supported by Cloudian. Cloudian can NOT be involved for any bugs or misconfiguration
due to this tool. So you are using it at your own risks and be aware of the restrictions.

--- RE-Synchronization from the bucket plaineco-rubrik-0 to the bucket plaineco-rubrik-0-tempo : ---
Do you want to continue (type : yes to continue,or anything else to abort) ? yes
OK, let's Go !
Transferred:          0 / 0 Bytes, -, 0 Bytes/s, ETA -
Checks:              92 / 92, 100%
Elapsed time:        600ms
Resync done.

Quick check if there are differences...
2020/01/29 18:57:08 NOTICE: S3 bucket plaineco-rubrik-0-tempo: 0 differences found
2020/01/29 18:57:08 NOTICE: S3 bucket plaineco-rubrik-0-tempo: 92 matching files
Checking the content of the bucket is possible by using the command :
./s3-bucket-migration.sh -c check -b plaineco-rubrik-0

Re-Sync is finished. You can adjust the protection policy on Cloudian - default policy will be selected
Please use reverse command. ./s3-bucket-migration.sh -c reverse -b plaineco-rubrik-0
[root@peter]#
```

Please, be sure that the Cloudian default protection policy is the new policy you want to apply for the new buckets. Example, before : RF3 and new : Erasure Coding 3+2

Optionnally, you can force a MD5 Hash on all objects before the reverse operation :

```
[root@peter]# ./s3-bucket-migration.sh -c check -b plaineco-rubrik-0

Disclaimer / Warning
Use this tool with precautions (review the config file created manually for a double-check) for your environment :
it is NOT an official tool supported by Cloudian. Cloudian can NOT be involved for any bugs or misconfiguration
due to this tool. So you are using it at your own risks and be aware of the restrictions.

--- Check the content of the buckets ---
Do you want to continue (type : yes to continue,or anything else to abort) ? yes
OK, let's Go !
MD5 calculations for plaineco-rubrik-0 and plaineco-rubrik-0-tempo
This might take long time and might use a lot of ressources ... depending on the structure of the buckets
Done. All checksums are in the log file : plaineco-rubrik-0.log

to continue with the process, you can adjust the protection policy on Cloudian - default policy will be selected
Then, use reverse command. ./s3-bucket-migration.sh -c reverse -b plaineco-rubrik-0
[root@peter]#
```

Purge the source bucket and remove it, then create a new one matching the same name :

```
[root@peter]# ./s3-bucket-migration.sh -c reverse -b plaineco-rubrik-0

Disclaimer / Warning
Use this tool with precautions (review the config file created manually for a double-check) for your environment :
it is NOT an official tool supported by Cloudian. Cloudian can NOT be involved for any bugs or misconfiguration
due to this tool. So you are using it at your own risks and be aware of the restrictions.

--- Reverse the operations from the bucket : plaineco-rubrik-0-tempo ---
Do you want to continue (type : yes to continue,or anything else to abort) ? yes
OK, let's Go !
Purging the bucket : plaineco-rubrik-0
Need to connect to the cluster by using SSH and root access...
root@s3-region.demo's password:
Bucket: plaineco-rubrik-0 purged.Creating a new bucket : plaineco-rubrik-0
Delete is finished, if there is no error, please Go Forward :-) ...
Then, use the command : ./s3-bucket-migration.sh -c copyback -b plaineco-rubrik-0
[root@peter]#
```

Copy all objects from the bucket “tempo” to the new bucket created with the new policy :

```
[root@peter]# ./s3-bucket-migration.sh -c copyback -b plaineco-rubrik-0

Disclaimer / Warning
Use this tool with precautions (review the config file created manually for a double-check) for your environment :
it is NOT an official tool supported by Cloudian. Cloudian can NOT be involved for any bugs or misconfiguration
due to this tool. So you are using it at your own risks and be aware of the restrictions.

--- Copy-back from the bucket : plaineco-rubrik-0-tempo to the bucket : plaineco-rubrik-0 ---
Do you want to continue (type : yes to continue,or anything else to abort) ? yes
OK, let's Go !
Copy-back in progress ...
Transferred:      13.841G / 13.841 GBytes, 100%, 54.451 MBytes/s, ETA 0s
Transferred:      92 / 92, 100%
Elapsed time:      4m20.2s

Quick check if there are differences...
2020/01/29 19:04:26 NOTICE: S3 bucket plaineco-rubrik-0-tempo: 0 differences found
2020/01/29 19:04:26 NOTICE: S3 bucket plaineco-rubrik-0-tempo: 92 matching files
Done.
Then, use the command : ./s3-bucket-migration.sh -c clean -b plaineco-rubrik-0
[root@peter]#
```

Clean the non-necessary objects and bucket (‘tempo’) :

```
[root@peter]# ./s3-bucket-migration.sh -c clean -b plaineco-rubrik-0

Disclaimer / Warning
Use this tool with precautions (review the config file created manually for a double-check) for your environment :
it is NOT an official tool supported by Cloudian. Cloudian can NOT be involved for any bugs or misconfiguration
due to this tool. So you are using it at your own risks and be aware of the restrictions.

--- Cleanup and deletion of non-necessary objects + bucket... ---
Do you want to continue (type : yes to continue,or anything else to abort) ? yes
OK, let's Go !
Purging the bucket : plaineco-rubrik-0-tempo
Need to connect to the cluster by using SSH and root access...
root@s3-region.demo's password:
Bucket: plaineco-rubrik-0-tempo purged.Job done. Nothing else to do.
Please resume the ArchivalLocation matching the current migration
[root@peter]#
```

Last step, you must resume the archival on the ArchivalLocation (Rubrik GUI) :

